

▼ Pima Indian Diabetes Prediction

The aim of this project to analyze the medical factors of a patient such as Glucose Level, Blood Pressure, Skin Thickness, Insulin Level and many others to predict whether the patient has diabetes or not.

About the Dataset This dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases. The objective of the dataset is to diagnostically predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset. Several constraints were placed on the selection of these instances from a larger database. In particular, all patients here are females at least 21 years old of Pima Indian heritage.

The datasets consists of several medical predictor variables and one target variable, Outcome. Predictor variables includes the number of pregnancies the patient has had, their BMI, insulin level, age, and so on.

Data Dictionary

Feature	Description
Pregnancies	Number of times pregnant
Glucose	Plasma glucose concentration a 2 hours in an oral glucose tolerance test
BloodPressure	Diastolic blood pressure (mm Hg)
SkinThickness	Triceps skin fold thickness (mm)
Insulin	2-Hour serum insulin (mu U/ml)
BMI	Body mass index (weight in kg/(height in m)^2)
DiabetesPedigreeFunction	Diabetes pedigree function
Age	Age (years)
Outcome	Class variable (0 or 1)

▼ Impact

The Pima Indian Diabetes Prediction project holds significant potential for impacting healthcare outcomes for Pima Indian females. Early detection and diagnosis of diabetes can play a crucial role in managing the condition effectively and preventing complications. By developing an accurate predictive model, healthcare providers can identify individuals at higher risk of diabetes and offer timely interventions and personalized treatment plans. This project's successful implementation may lead to improved health management strategies, reduced healthcare costs, and enhanced overall well-being for the Pima Indian female community. Additionally, the insights gained from this study may contribute to broader research on diabetes risk factors and aid in formulating targeted public health initiatives for diabetes prevention and awareness within the Pima Indian population. The ethical and responsible use of data in this project will be ensured to protect patient privacy and promote transparency in the application of predictive modeling in healthcare settings.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

df=pd.read_csv('diabetes.csv')

df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

```
df.tail()
```

```
df.info()
```

```
df.shape
```

(768, 9)

```
df.describe()
```

```
df.nunique()
```

Checking the unique values for each variable in the dataset

```
#checking unique values
```

```
variables = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome']
for i in variables:
    print(df[i].unique())
```

[6	1	8	0	5	3	10	2	4	7	9	11	13	15	17	12	14]						
[148	85	183	89	137	116	78	115	197	125	110	168	139	189	166	100	118	107						
103	126	99	196	119	143	147	97	145	117	109	158	88	92	122	138	102	90						
111	180	133	106	171	159	146	71	105	101	176	150	73	187	84	141	114							
95	129	79	0	62	131	112	113	74	83	136	80	123	81	134	142	144	93						
163	151	96	155	76	160	124	162	132	120	173	170	128	108	154	57	156	153						
188	152	104	87	75	179	130	194	181	135	184	140	177	164	91	165	86	193						
191	161	167	77	182	157	178	61	98	127	82	72	172	94	175	195	68	186						
198	121	67	174	199	56	169	149	65	190]														
[72	66	64	40	74	50	0	70	96	92	80	60	84	30	88	90	94	76					
82	75	58	78	68	110	56	62	85	86	48	44	65	108	55	122	54	52						
98	104	95	46	102	100	61	24	38	106	114]													
[35	29	0	23	32	45	19	47	38	30	41	33	26	15	36	11	31	37	42	25	18	24	39	27
21	34	10	60	13	20	22	28	54	40	51	56	14	17	50	44	12	46	16	7	52	43	48	8
49	63	99]																					
[0	94	168	88	543	846	175	230	83	96	235	146	115	140	110	245	54	192					
207	70	240	82	36	23	300	342	304	142	128	38	100	90	270	71	125	176						
48	64	228	76	220	40	152	18	135	495	37	51	99	145	225	49	50	92						
325	63	284	119	204	155	485	53	114	105	285	156	78	130	55	58	160	210						

```

318 44 190 280 87 271 129 120 478 56 32 744 370 45 194 680 402 258
375 150 67 57 116 278 122 545 75 74 182 360 215 184 42 132 148 180
205 85 231 29 68 52 255 171 73 108 43 167 249 293 66 465 89 158
84 72 59 81 196 415 275 165 579 310 61 474 170 277 60 14 95 237
191 328 250 480 265 193 79 86 326 188 106 65 166 274 77 126 330 600
185 25 41 272 321 144 15 183 91 46 440 159 540 200 335 387 22 291
392 178 127 510 16 112]
[33.6 26.6 23.3 28.1 43.1 25.6 31. 35.3 30.5 0. 37.6 38. 27.1 30.1
25.8 30. 45.8 29.6 43.3 34.6 39.3 35.4 39.8 29. 36.6 31.1 39.4 23.2
22.2 34.1 36. 31.6 24.8 19.9 27.6 24. 33.2 32.9 38.2 37.1 34. 40.2
22.7 45.4 27.4 42. 29.7 28. 39.1 19.4 24.2 24.4 33.7 34.7 23. 37.7
46.8 40.5 41.5 25. 25.4 32.8 32.5 42.7 19.6 28.9 28.6 43.4 35.1 32.
24.7 32.6 43.2 22.4 29.3 24.6 48.8 32.4 38.5 26.5 19.1 46.7 23.8 33.9
20.4 28.7 49.7 39. 26.1 22.5 39.6 29.5 34.3 37.4 33.3 31.2 28.2 53.2
34.2 26.8 55. 42.9 34.5 27.9 38.3 21.1 33.8 30.8 36.9 39.5 27.3 21.9
40.6 47.9 50. 25.2 40.9 37.2 44.2 29.9 31.9 28.4 43.5 32.7 67.1 45.
34.9 27.7 35.9 22.6 33.1 30.4 52.3 24.3 22.9 34.8 30.9 40.1 23.9 37.5
35.5 42.8 42.6 41.8 35.8 37.8 28.8 23.6 35.7 36.7 45.2 44. 46.2 35.
43.6 44.1 18.4 29.2 25.9 32.1 36.3 40. 25.1 27.5 45.6 27.8 24.9 25.3
37.9 27. 26. 38.7 20.8 36.1 30.7 32.3 52.9 21. 39.7 25.5 26.2 19.3
38.1 23.5 45.5 23.1 39.9 36.8 21.8 41. 42.2 34.4 27.2 36.5 29.8 39.2
38.4 36.2 48.3 20. 22.3 45.7 23.7 22.1 42.1 42.4 18.2 26.4 45.3 37.
24.5 32.2 59.4 21.2 26.7 30.2 46.1 41.3 38.8 35.2 42.3 40.7 46.5 33.5
37.3 30.3 26.3 21.7 36.4 28.5 26.9 38.6 31.3 19.5 20.1 40.8 23.4 28.3
38.9 57.3 35.6 49.6 44.6 24.1 44.5 41.2 49.3 46.3]
[0.627 0.351 0.672 0.167 2.288 0.201 0.248 0.134 0.158 0.232 0.191 0.537
1.441 0.398 0.587 0.484 0.551 0.254 0.183 0.529 0.704 0.388 0.451 0.263
0.205 0.257 0.487 0.245 0.337 0.546 0.851 0.267 0.188 0.512 0.966 0.42
0.665 0.503 1.39 0.271 0.696 0.235 0.721 0.294 1.893 0.564 0.586 0.344
0.305 0.491 0.526 0.342 0.467 0.718 0.962 1.781 0.173 0.304 0.27 0.699
0.258 0.203 0.855 0.845 0.334 0.189 0.867 0.411 0.583 0.231 0.396 0.14
0.391 0.37 0.307 0.102 0.767 0.237 0.227 0.698 0.178 0.324 0.153 0.165
0.443 0.261 0.277 0.761 0.255 0.13 0.323 0.356 0.325 1.222 0.179 0.262
0.283 0.93 0.801 0.207 0.287 0.336 0.247 0.199 0.543 0.192 0.588 0.539
0.22 0.654 0.223 0.759 0.26 0.404 0.186 0.278 0.496 0.452 0.403 0.741
0.361 1.114 0.457 0.647 0.088 0.597 0.532 0.703 0.159 0.268 0.286 0.318
0.272 0.572 0.096 1.4 0.218 0.085 0.399 0.432 1.189 0.687 0.137 0.637
0.833 0.229 0.817 0.204 0.368 0.743 0.722 0.256 0.709 0.471 0.495 0.18
0.542 0.773 0.678 0.719 0.382 0.319 0.19 0.956 0.084 0.725 0.299 0.244

```

In the dataset the variables except Pregnancies and Outcome cannot have value as 0, because it is not possible to have 0 Glucose Level or to have 0 Blood Pressure. So, this will be counted as incorrect information

Checking the count of value 0 in the variables

```

variables = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age',]
for i in variables:
    c = 0
    for x in (df[i]):
        if x == 0:
            c = c + 1
    print(i,c)

Glucose 5
BloodPressure 35
SkinThickness 227
Insulin 374
BMI 11
DiabetesPedigreeFunction 0
Age 0

```

Replacing the 0 value in the variables - Glucose, BloodPressure, SkinThickness, Insulin, BMI

```

# Fill missing values in each column with the mean of that column
variables = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']
for column in variables:
    mean_value = df[column].mean()
    df[column].fillna(mean_value, inplace=True)

missing_value_counts = {}

# Loop through columns and count missing values
for column in variables:
    missing_count = df[column].isnull().sum()
    missing_value_counts[column] = missing_count


print(missing_value_counts)

{'Glucose': 0, 'BloodPressure': 0, 'SkinThickness': 0, 'Insulin': 0, 'BMI': 0}

```

▼ Exploratory Data Analysis





```
plt.figure(figsize=(5,5))
plt.pie(df['Outcome'].value_counts(), labels=['No Diabetes', 'Diabetes'], autopct='%1.1f%%', shadow=False, startangle=90)
plt.title('Diabetes Outcome')
plt.show()
```


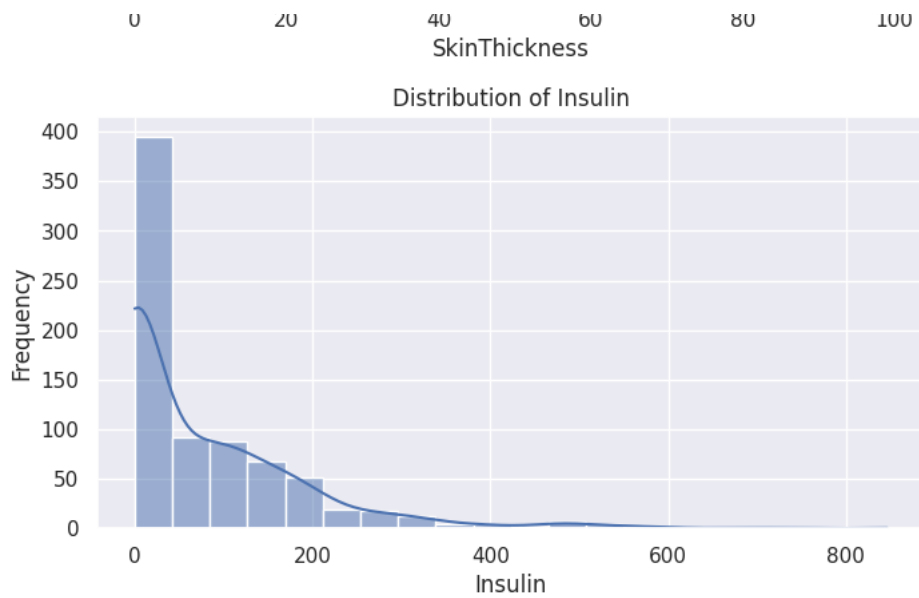


```
numeric_columns = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age']
summary_stats = df[numeric_columns].describe()
```

```
# Distribution plots for numeric columns
```

```
for column in numeric_columns:
    plt.figure(figsize=(8, 4))
    sns.histplot(df[column], kde=True, bins=20)
    plt.title(f'Distribution of {column}')
    plt.xlabel(column)
    plt.ylabel('Frequency')
    plt.show()
```






```
# Distribution of the 'Outcome' column (assuming it's binary)
plt.figure(figsize=(6, 4))
sns.countplot(data=df, x='Outcome')
plt.title('Distribution of Outcome')
plt.xlabel('Outcome')
plt.ylabel('Count')
plt.show()
```

Distribution of Outcome

Age Distribution and Diabetes

```
sns.catplot(x="Outcome", y="Age", kind="swarm", data=df)
```

```
/usr/local/lib/python3.10/dist-packages/seaborn/categorical.py:3544: UserWarning: 7.2% of the points cannot be placed; you may want warnings.warn(msg, UserWarning)
<seaborn.axisgrid.FacetGrid at 0x798aecc54d00>/usr/local/lib/python3.10/dist-packages/seaborn/categorical.py:3544: UserWarning: 23.2% of the points cannot be plac
warnings.warn(msg, UserWarning)
```





From the graph, it is quite clear that majority of the patients are adult within the age group of 20-30 years. Patients in the age range 40-55 years are more prone to diabetes, as compared to other age groups. Since the number adults in the age group 20-30 years is more, the number of patients with diabetes is also more as compared of other age groups.

Pregnancies and Diabetes

```
fig,ax = plt.subplots(1,2,figsize=(15,5))
sns.boxplot(x='Outcome',y='Pregnancies',data=df,ax=ax[0])
sns.violinplot(x='Outcome',y='Pregnancies',data=df,ax=ax[1])
```

```
<Axes: xlabel='Outcome', ylabel='Pregnancies'>
```




Both boxplot and violinplot shows strange relation between the number of pregnancies and diabetes. According to the graphs the increased number of pregnancies highlights increased risk of diabetes.

Glucose and Diabetes

```
sns.boxplot(x='Outcome', y='Glucose', data=df).set_title('Glucose vs Diabetes')
```

```
Text(0.5, 1.0, 'Glucose vs Diabetes')
```

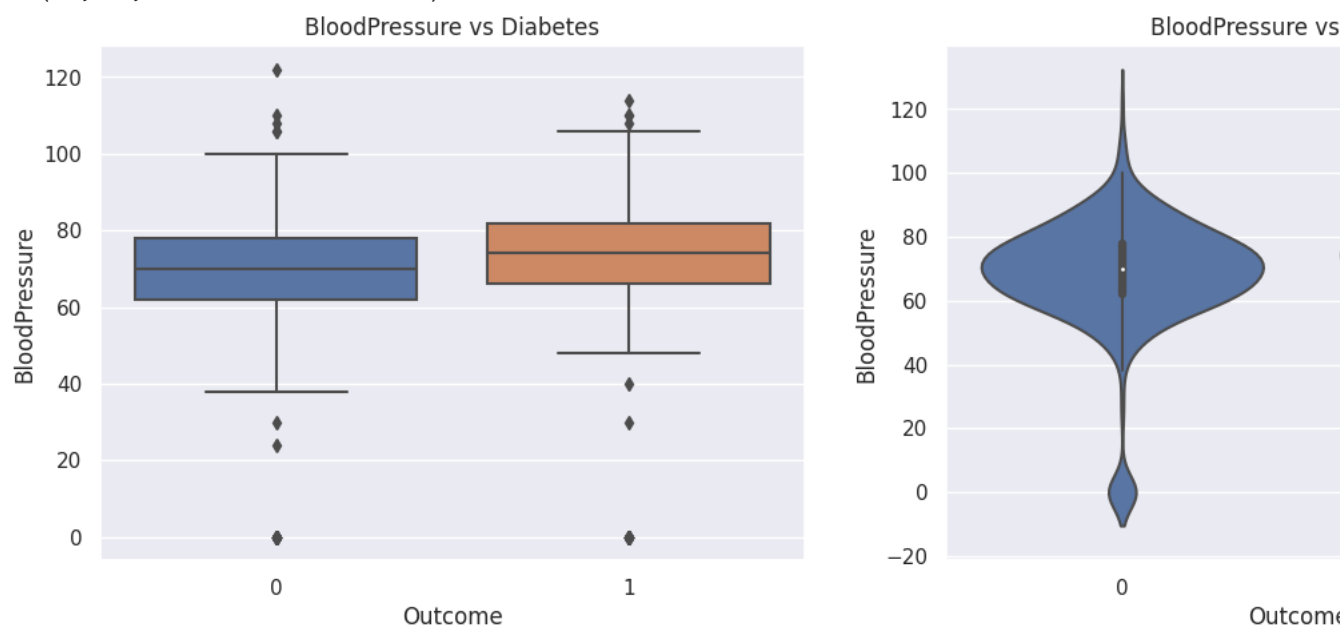


Glucose level plays a major role in determine whether the patient is diabetic or not. The patients with median glucocse level less than 120 are more likely to be non-diabetic. The patients with median glucocse level greather than 140 are more likely to be diabetic. Therefore, high glucose levels is a good indicator of diabetes.

Blood Pressuse and Diabetes

```
fig,ax = plt.subplots(1,2,figsize=(15,5))
sns.boxplot(x='Outcome', y='BloodPressure', data=df, ax=ax[0]).set_title('BloodPressure vs Diabetes')
sns.violinplot(x='Outcome', y='BloodPressure', data=df, ax=ax[1]).set_title('BloodPressure vs Diabetes')
```

```
Text(0.5, 1.0, 'BloodPressure vs Diabetes')
```




Both the boxplot and voilinplot provides clear understanding of the realtion between the blood pressure and diabetes. The boxplot shows that the median of the blood pressure for the diabetic patients is slightly higher than the non-diabetic patients. The voilinplot shows that the

distribution of the blood pressure for the diabetic patients is slightly higher than the non-diabetic patients. But there has been not enough evidence to conclude that the blood pressure is a good predictor of diabetes.

Skin Thickness and Diabetes

```
fig,ax = plt.subplots(1,2,figsize=(15,5))
sns.boxplot(x='Outcome', y='SkinThickness', data=df,ax=ax[0]).set_title('SkinThickness vs Diabetes')
sns.violinplot(x='Outcome', y='SkinThickness', data=df,ax=ax[1]).set_title('SkinThickness vs Diabetes')
```

Text(0.5, 1.0, 'SkinThickness vs Diabetes')



Here both the boxplot and violinplot reveals the effect of diabetes on skin thickness. As observed in the boxplot, the median of skin thickness is higher for the diabetic patients than the non-diabetic patients, where non-diabetic patients have median skin thickness near 20 in comparison to skin thickness nearly 30 in diabetic patients. The violinplot shows the distribution of patients' skin thickness among the patients, where the non-diabetic ones have greater distribution near 20 and diabetic much less distribution near 20 and increased distribution near 30. Therefore, skin thickness can be an indicator of diabetes.

Insulin and Diabetes


```
fig,ax = plt.subplots(1,2,figsize=(15,5))
sns.boxplot(x='Outcome',y='Insulin',data=df,ax=ax[0]).set_title('Insulin vs Diabetes')
sns.violinplot(x='Outcome',y='Insulin',data=df,ax=ax[1]).set_title('Insulin vs Diabetes')
```

Insulin is a major body hormone that regulates glucose metabolism. Insulin is required for the body to efficiently use sugars, fats and proteins. Any change in insulin amount in the body would result in change glucose levels as well. Here the boxplot and violinplot shows the distribution of insulin level in patients. In non diabetic patients the insulin level is near to 100, whereas in diabetic patients the insulin level is near to 200. In the violinplot we can see that the distribution of insulin level in non diabetic patients is more spread out near 100, whereas in diabetic patients the distribution is contracted and shows a little bit spread in higher insulin levels. This shows that the insulin level is a good indicator of diabetes.

BMI and Diabetes

```
fig,ax = plt.subplots(1,2,figsize=(15,5))
sns.boxplot(x='Outcome',y='BMI',data=df,ax=ax[0])
sns.violinplot(x='Outcome',y='BMI',data=df,ax=ax[1])
```

<Axes: xlabel='Outcome', ylabel='BMI'>



Both graphs highlights the role of BMI in diabetes prediction. Non diabetic patients have a normal BMI within the range of 25-35 whereas the diabetic patients have a BMI greater than 35. The violinplot reveals the BMI distribution, where the non diabetic patients have a increased spread from 25 to 35 with narrows after 35. However in diabetic patients there is increased spread at 35 and increased spread 45-50 as compared to non diabetic patients. Therefore BMI is a good predictor of diabetes and obese people are more likely to be diabetic.

Diabetes Pedigree Function and Diabetes Outcome

```
fig,ax = plt.subplots(1,2,figsize=(15,5))
sns.boxplot(x='Outcome',y='DiabetesPedigreeFunction',data=df,ax=ax[0]).set_title('Diabetes Pedigree Function')
sns.violinplot(x='Outcome',y='DiabetesPedigreeFunction',data=df,ax=ax[1]).set_title('Diabetes Pedigree Function')
```

```
Text(0.5, 1.0, 'Diabetes Pedigree Function')
```

Diabetes Pedigree Function

Diabetes Pedigree

2.5

Diabetes Pedigree Function (DPF) calculates diabetes likelihood depending on the subject's age and his/her diabetic family history. From the boxplot, the patients with lower DPF, are much less likely to have diabetes. The patients with higher DPF, are much more likely to have diabetes. In the violinplot, majority of the non diabetic patients have a DPF of 0.25-0.35, whereas the diabetic patients have a increased DPF, which is shown by the their distribution in the violinplot where there is a increased spread in the DPF from 0.5 -1.5. Therefore the DPF is a good indicator of diabetes.

%


%

Coorelation Matrix Heatmap

%

%

```
# 1. Categorical vs. Categorical
plt.figure(figsize=(8, 6))
sns.countplot(data=df, x='Pregnancies', hue='Outcome')
plt.title('Distribution of Pregnancies by Outcome')
plt.xlabel('Pregnancies')
plt.ylabel('Count')
plt.legend(title='Outcome', labels=['No Diabetes', 'Diabetes'])
plt.show()
```



```
from scipy.stats import chi2_contingency, ttest_ind
# Chi-squared test for independence
contingency_table = pd.crosstab(df['Pregnancies'], df['Outcome'])
chi2, p, _, _ = chi2_contingency(contingency_table)
print(f"Chi-squared p-value: {p}")
```

Chi-squared p-value: 8.648349123362548e-08

The p-value from a chi-squared test for independence represents the probability of observing the observed association (or more extreme) between two categorical variables, such as 'Pregnancies' and 'Outcome,' under the assumption that there is no actual association (independence) between them.


where the p-value is approximately 8.65e-08 (very close to zero), this extremely small p-value indicates the following:

Statistical Significance: The chi-squared test has found strong evidence to reject the null hypothesis. The null hypothesis in this context would typically be that there is no association or independence between 'Pregnancies' and 'Outcome.'

Association or Dependence: The small p-value suggests that there is a significant association or dependence between the number of pregnancies ('Pregnancies') and the outcome of the diabetes test ('Outcome') in your dataset.

```
# 2. Categorical vs. Numerical
plt.figure(figsize=(8, 6))
sns.boxplot(data=df, x='Outcome', y='Age')
plt.title('Age vs. Outcome')
```

```
plt.xlabel('Outcome')
plt.ylabel('Age')
plt.show()
```



```
# t-test for difference in means
no_diabetes_age = df[df['Outcome'] == 0]['Age']
diabetes_age = df[df['Outcome'] == 1]['Age']
t_stat, p_value = ttest_ind(no_diabetes_age, diabetes_age)
print(f"t-test p-value for Age and Outcome: {p_value}")


t-test p-value for Age and Outcome: 2.2099754606654358e-11
```

The p-value you obtained from the t-test for 'Age' and 'Outcome' is approximately 2.21×10^{-11} , which is an extremely small number. In simple terms, here's what you can understand from this result:

Statistical Significance: The small p-value (much less than 0.05 or any common significance level) indicates strong evidence against the idea that there is no difference in the average ages between the two groups (diabetes and no diabetes).

Difference in Means: A small p-value suggests that there is a statistically significant difference in the average ages of individuals with diabetes and those without diabetes.

```
# 3. Numerical vs. Numerical
plt.figure(figsize=(8, 6))
sns.scatterplot(data=df, x='BMI', y='BloodPressure', hue='Outcome')
plt.title('BMI vs. BloodPressure by Outcome')
plt.xlabel('BMI')
plt.ylabel('BloodPressure')
plt.show()
```



```
#Correlation between BMI and BloodPressure
correlation = df[['BMI', 'BloodPressure']].corr(method='pearson')
print("Correlation matrix between BMI and BloodPressure:")
print(correlation)
```

```
Correlation matrix between BMI and BloodPressure:
          BMI  BloodPressure
BMI      1.000000      0.281805
BloodPressure 0.281805      1.000000
```

Correlation Coefficient Values:

The correlation coefficient between 'BMI' and 'BloodPressure' is approximately 0.282. This value is between -1 and 1. A positive correlation coefficient (0.282) indicates a positive linear relationship, meaning that as 'BMI' increases, 'BloodPressure' tends to increase as well. However, the correlation is relatively weak, as the coefficient is closer to 0 than to 1. Strength of Relationship:

The correlation coefficient of 0.282 suggests a relatively weak positive relationship between 'BMI' and 'BloodPressure.' This means that while there is a tendency for these variables to increase together, the relationship is not very strong. Interpretation:

In practical terms, this correlation suggests that there is some degree of association between higher BMI values and higher blood pressure values. This is consistent with the common understanding that excess body weight (higher BMI) can be a risk factor for elevated blood pressure. Magnitude of Correlation:

The magnitude of the correlation coefficient (0.282) indicates that while there is a relationship, it may not be the only factor influencing blood pressure. Other factors could also play a significant role.

In summary, the correlation matrix suggests a weak positive linear relationship between 'BMI' and 'BloodPressure' in your dataset. It's important to note that correlation does not imply causation. While these variables are correlated, it doesn't necessarily mean that one directly causes the other; other factors and mechanisms may be involved.


```
correlation_matrix=df.corr()

plt.figure(figsize=(10, 8)) # Adjust the figure size as needed

# Create a heatmap using Seaborn
sns.set(font_scale=1.0) # Adjust the font size as needed
sns.heatmap(correlation_matrix, annot=True, fmt=".2f", cmap="coolwarm", linewidths=0.5)

# Customize the plot
plt.title("Correlation Heatmap")
plt.xticks(rotation=45)
plt.yticks(rotation=45)

# Display the heatmap
plt.show()
```



```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df.drop('Outcome',axis=1),df['Outcome'],test_size=0.2,random_state=42)
```

Logistic Regression

```
#building model
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr
```

```
#training the model
lr.fit(X_train,y_train)
#training accuracy
lr.score(X_train,y_train)
```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status STOP: TOTAL NO. of ITERATIONS REACHED LIMIT).

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
0.7719869706840391
```

```
#predicted outcomes
lr_pred = lr.predict(X_test)
```

```
from sklearn.metrics import confusion_matrix
sns.heatmap(confusion_matrix(y_test, lr_pred), annot=True, cmap='Blues')
plt.xlabel('Predicted Values')
plt.ylabel('Actual Values')
plt.title('Confusion Matrix for Logistic Regression')
plt.show()
```

Confusion Matrix for Logistic Regression



```
ax = sns.distplot(y_test, color='r', label='Actual Value',hist=False)
sns.distplot(lr_pred, color='b', label='Predicted Value',hist=False,ax=ax)
plt.title('Actual vs Predicted Value Logistic Regression')
plt.xlabel('Outcome')
plt.ylabel('Count')
```

<ipython-input-48-bc104d462945>:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
ax = sns.distplot(y_test, color='r', label='Actual Value',hist=False)
```

<ipython-input-48-bc104d462945>:2: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(lr_pred, color='b', label='Predicted Value',hist=False,ax=ax)
Text(0, 0.5, 'Count')
```



These distribution plot clearly visualizes the accuracy of the model. The red color represents the actual values and the blue color represents the predicted values. The more the overlapping of the two colors, the more accurate the model is.

```
from sklearn.metrics import classification_report
print(classification_report(y_test, lr_pred))
```

	precision	recall	f1-score	support
0	0.81	0.79	0.80	99
1	0.64	0.67	0.65	55
accuracy			0.75	154
macro avg	0.73	0.73	0.73	154
weighted avg	0.75	0.75	0.75	154

▼ Random Forest Classifier

```

#buidling model
from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier(n_estimators=100,random_state=42)
rfc


#training model
rfc.fit(X_train, y_train)
#training accuracy
rfc.score(X_train, y_train)

1.0

#predicted outcomes
rfc_pred = rfc.predict(X_test)

sns.heatmap(confusion_matrix(y_test, rfc_pred), annot=True, cmap='Blues')
plt.xlabel('Predicted Values')
plt.ylabel('Actual Values')
plt.title('Confusion Matrix for Logistic Regression')
plt.show()

```



```

ax = sns.distplot(y_test, color='r', label='Actual Value',hist=False)
sns.distplot(rfc_pred, color='b', label='Predicted Value',hist=False,ax=ax)
plt.title('Actual vs Predicted Value Logistic Regression')
plt.xlabel('Outcome')
plt.ylabel('Count')

```



```
<ipython-input-53-9669e741e5cd>:1: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
ax = sns.distplot(y_test, color='r', label='Actual Value', hist=False)
```

```
<ipython-input-53-9669e741e5cd>:2: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
print(classification_report(y_test, rfc_pred))
```

	precision	recall	f1-score	support
0	0.79	0.78	0.78	99
1	0.61	0.62	0.61	55
accuracy			0.72	154
macro avg	0.70	0.70	0.70	154
weighted avg	0.72	0.72	0.72	154

```
100
```

▼ Support Vector Machine (SVM)

```
8
```

```
#building model
```

```
from sklearn.svm import SVC
```

```
svm = SVC(kernel='linear', random_state=0)
```

```
svm
```

```
#training the model
```

```
svm.fit(X_train, y_train)
```

```
#training the model
```

```
svm.score(X_test, y_test)
```

```
0.7532467532467533
```

```
#predicting outcomes
```

```
svm_pred = svm.predict(X_test)
```


```
sns.heatmap(confusion_matrix(y_test, svm_pred), annot=True, cmap='Blues')
```

```
plt.xlabel('Predicted Values')
```

```
plt.ylabel('Actual Values')
```

```
plt.title('Confusion Matrix for Logistic Regression')
```

```
plt.show()
```



```
ax = sns.distplot(y_test, color='r', label='Actual Value',hist=False)
sns.distplot(svm_pred, color='b', label='Predicted Value',hist=False,ax=ax)
plt.title('Actual vs Predicted Value Logistic Regression')
plt.xlabel('Outcome')
plt.ylabel('Count')
```

<ipython-input-59-b9a6ee476682>:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
ax = sns.distplot(y_test, color='r', label='Actual Value',hist=False)
```

<ipython-input-59-b9a6ee476682>:2: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(svm_pred, color='b', label='Predicted Value',hist=False,ax=ax)
Text(0, 0.5, 'Count')
```



```
print(classification_report(y_test, rfc_pred))
```

	precision	recall	f1-score	support
0	0.79	0.78	0.78	99
1	0.61	0.62	0.61	55
accuracy			0.72	154
macro avg	0.70	0.70	0.70	154
weighted avg	0.72	0.72	0.72	154

▼ Comparing the models

```
#comparing the accuracy of different models
```

```
sns.barplot(x=['Logistic Regression', 'RandomForestClassifier', 'SVM'], y=[0.7792207792207793,0.7662337662337663,0.7597402597402597])
```

```
plt.xlabel('Classifier Models')
```

```
plt.ylabel('Accuracy')
```

```
plt.title('Comparison of different models')
```